

# Lecture 3: Introduction to PINNs and the DRM

**Chris Budd and Aengus Roberts<sup>1</sup>**

<sup>1</sup>University of Bath

Ai4Sci, December 2025

## Some papers to look at

- Raissi, M., P. Perdikaris, and G. E. Karniadakis. 2019. "Physics-Informed Neural Networks: A Deep Learning Framework for Solving Forward and Inverse Problems Involving Nonlinear Partial Differential Equations." *Journal of Computational Physics* 378: 686–707. <https://doi.org/https://doi.org/10.1016/j.jcp.2018.10.045>.
- Grossman et. al. *Can PINNS beat the FE method?*
- Shin et. al. *On the convergence of PINNS for linear elliptic and parabolic PDEs*
- Wang et. al. *When and why PINNS fail to train: a neural tangent kernel perspective*
- Kiyani *Which Optimizer Works Best for Physics-Informed Neural Networks and Kolmogorov-Arnold Networks?*
- Chen et. al. *Sidecar: A Structure-Preserving Framework for Solving Partial Differential Equations with Neural Networks*
- Russell et. al. *Two-point boundary value problems*

# Motivation: Solving ODEs and PDEs

Seek to solve ODE/PDE problems of the form

$$\mathbf{u}_t = F(\mathbf{x}, u, \nabla u, \nabla^2 u) \quad \text{with BC}$$

eg. ODE (IVP or BVP)

$$\frac{du}{dt} = u^2, \quad u(0) = 1, \quad -\epsilon^2 \frac{d^2 u}{dx^2} + u = 1, \quad u(0) = u(1) = 0.$$

eg. PDE

$$-\Delta u = f(x), \quad iu_t + \Delta u + u|u|^2 = 0, \quad u_t = \Delta u + f(x, u) \quad x \in R^n$$

# Classical Method 1. Finite Differences

Work with a set of **point values** for a function **never with a function directly**  
IVP/BVP:

$$U_j^n \approx u(n\Delta t, j\Delta x), \quad \Delta t, \Delta x \ll 1$$

eg.

$$u_t = u_{xx} + u^3$$

IMEX Crank-Nicholson method:

$$\frac{U_j^{n+1} - U_j^n}{\Delta t} =$$

$$\frac{1}{2\Delta x^2} \left( U_{j+1}^{n+1} - 2U_j^{n+1} + U_{j-1}^{n+1} + U_{j+1}^n - 2U_j^n + U_{j-1}^n \right) + (U_j^n)^3.$$

- Have error estimates of the form

$$\|U_j^n - u(n\Delta t, j\Delta x)\| < C(u)\Delta t^p \Delta x^q, \quad \Delta t, \Delta x \rightarrow 0$$

- Can reduce  $C(u)$  and increase  $p, q$  using an adaptive approach.
- Lots of software
- Have to recover the function from the point values
- Awkward in higher dimensions!

## Classical Method 2: Finite Elements

Express  $u(x, t)$  as a Galerkin approximation:

$$u(t, x) \approx U(t, x) = \sum_{i=0}^N U_i(t) \phi_i(x)$$

with  $\phi_i(x)$  **Not** globally differentiable *locally supported, piece-wise polynomial spline functions*

- Work with a function  $U(x) \in H^1$
- Require  $U$  to satisfy a *weak form* of the PDE (see Lecture 4)
- Have guaranteed error estimates of the form

$$\|u - U\|_{H^1} < C(u)N^{-\alpha}$$

- Lots of software (see Lecture 2b)
- Can reduce  $C(u)$  and increase  $\alpha$  using an adaptive approach.
- Awkward in higher dimensions!

## Classical Method 3: Collocation

Express  $u(x, t)$  as a function approximation:

$$u(t, x) \approx U(t, x) = \sum_{i=0}^N U_i(t) \phi_i(x)$$

with  $\phi_i(x)$ :

**Locally differentiable locally supported, piece-wise polynomial spline functions such as Lagrange polynomial interpolants**

- Work with a function  $U(x) \in C_{loc}^n$
- Require  $U$  to satisfy the PDE **exactly** at a carefully chosen set of **collocation points**

- Have guaranteed error estimates of the form

$$\|u - U\|_{C_2} < C(u)N^{-\alpha}$$

- Can reduce  $C(u)$  and increase  $\alpha$  using VERY carefully chosen collocation points
- Implemented in python as `solve_bvp`, Colsys
- Impossible in higher dimensions!

# Issues with classical methods


- Well tested and good convergence theories
- Accuracy of the computation depends crucially on the choice and shape of the mesh points
- Generally require **time and practice** to implement (although FireDrake etc. makes this easier)
- Two-fold curse of dimensionality

# PINNS

Physics Informed Neural Networks for solving PDEs: advertised as "Mesh free methods".

Use a Deep Neural Net of width  $W$  and depth  $L$  to give a **nonlinear functional approximation** to  $u(\mathbf{x})$  with **input**  $\mathbf{x}$ .

$$y(\mathbf{x}) = DNN(\mathbf{x})$$



$y(\mathbf{x})$  is constructed via a combination of linear transformations and nonlinear/semi-linear activation functions.

Example: Shallow 1D neural net

$$y(\mathbf{x}) = \sum_{i=0}^{W-1} c_i \sigma(a_i \mathbf{x} + \mathbf{b}_i)$$

Often take

$$\sigma(z) = \text{ReLU}^3(z) \equiv z_+^3, \quad \text{or} \quad \sigma(z) = \tanh(z)$$

As a result we can define  $y_{zz}$  for every point  $z$

# Operation of a 'traditional' PINN

Want to solve a ODE/PDE in  $\mathbf{x}$

- PINNs are similar to collocation methods
- Assume that  $y(\mathbf{x})$  has strong regularity eg.  $C^2$
- Differentiate  $y(\mathbf{x})$  exactly using the chain rule
- Evaluate the **PDE residual** at  $N_r$  collocation points  $\mathbf{X}_i$ , chosen to be uniformly spaced, or samples from a random distribution
- Train the neural net to minimise a **loss function**  $L$  combining the PDE residual and boundary and initial conditions. May use **Epochs linked to the random training samples.**

## Eg 1. Solution of regular two-point BVPs by PINNs

Consider the two-point BVP with Dirichlet boundary conditions:

$$-u_{xx} = f(x, u, u_x), \quad x \in [0, 1] \quad u(0) = a, \quad u(1) = b.$$

Define output of the PINN by  $y(x)$  and residual  $r(x) := y_{xx} + f(x, y, y_x)$ .  
Train the coefficients  $\theta$  of the NN by minimising the loss function

$$L = \frac{1}{N_r} \sum_i^{N_r} |r(X_i^r)|^2 + \frac{\beta}{2} (|y(0) - a|^2 + |y(1) - b|^2),$$

where  $\{X_i^r\}_i^{N_r}$  are the **collocation points** (possibly randomly) placed in  $(0, 1)$ .

## Eg 2. Solution of regular second order IVPs by PINNs

Consider the second order IVP with initial conditions:

$$-u_{xx} = f(x, u, u_x), \quad x \in [0, 1] \quad u(0) = a, \quad u'(0) = b.$$

Define output of the PINN by  $y(x)$  and residual  $r(x) := y_{xx} + f(x, y, y_x)$ .  
Train the coefficients  $\theta$  of the NN by minimising the loss function

$$L = \frac{1}{N_r} \sum_i^{N_r} |r(X_i^r)|^2 + \frac{\beta}{2} (|y(0) - a|^2 + |y'(0) - b|^2),$$

where  $\{X_i^r\}_i^{N_r}$  are the **collocation points** (possibly randomly) placed in  $(0, 1)$ .

## Eg 3. Solution of regular parabolic PDEs by time-stepping PINNs

Consider the **semilinear parabolic PDE** with Dirichlet boundary conditions:

$$u_t = u_{xx} + f(x, u, u_x), \quad x \in [0, 1] \quad u(0) = a, \quad u(1) = b$$

and initial condition

$$u(x, 0) = u_0(x).$$

Implicit time-stepping scheme  $U^n(x) \approx u(n\Delta t, x)$ .

$$\frac{U^{n+1}(x) - U^n(x)}{\Delta t} = U_{xx}^{n+1} + f(x, U^{n+1}, U_x^{n+1}) \equiv F(U^{n+1}).$$

- Start with  $U^0(x) = u_0(x)$
- For  $n > 0$  : Define output of the PINN by  $y(x)$  and residual

$$r(x) := U^n(x) + \Delta t F(y).$$

- The PINN is trained by minimising the loss function

$$L = \frac{1}{N_r} \sum_i^{N_r} |r(X_i^r)|^2 + \frac{1}{2} (|y(0) - a|^2 + |y(1) - b|^2),$$

where  $\{X_i^r\}_i^{N_r}$  are the **collocation points** placed in  $(0, 1)$ .

- Set  $U^{n+1} = y(x)$  and repeat.

## Eg 4. Solution of regular parabolic PDEs by a full PINN

Consider the **same semilinear parabolic PDE**

$$u_t = u_{xx} + f(x, u, u_x), \quad x \in [0, 1] \quad u(0) = a, \quad u(1) = b.$$

**Bold approach!**

- Have NN with  $y(t, x)$  a function of  $t$  and  $x$
- Take  $N_r$  collocation points  $Z_i$  in **space and time**
- Residual  $r = u_t - u_{xx} - f(x, u, u_x)$
- Minimise

$$\begin{aligned} L = & \frac{1}{N_r} \sum_i^{N_r} |r(Z_i^r)|^2 + \frac{\beta}{2} (|y(0) - a|^2 + |y(1) - b|^2) \\ & + \gamma \|y(0, x) - u_0(x)\|^2. \end{aligned}$$

**Problematic as space and time play very different roles in the PDE**

Numerical results for:  $-u'' = \pi^2 \sin(\pi x)$ .





Figure: Residual based Loss and  $L^2$  error of the PINN solution for  $N_r = 100$

# Numerical results: $u(x) = \sin(\pi x)$



**Figure:** (Left) PINN with depth  $L = 2$  and width  $W = 30$ .  $N_r = 100$  uniformly distributed quadrature points, activation function:  $\tanh$ , optimizer: Adam with  $lr = 1e-3$  (Right) Convergence rate for 1st order interpolant

## Eg 2. Singular Reaction-Diffusion Equation

Solve  $-\varepsilon^2 u_{xx} + u = 1 - x$  on  $[0, 1]$     $u(0) = u(1) = 0$






Figure: PINN (tanh) trained for 20000 epochs,  $N_r = 101$ , Adam optimizer with  $\text{lr} = 1\text{e}-3$ . (left) Uniform collocation points (right) Adapted collocation points **much faster training**

## Eg 3. Bad news: Convection-dominated equation



PINNs often fail to train when the solution of the BVP exhibits singular and convective behaviour [Krishnapriyan, Aditi et. al., (2021)]:

$$-\varepsilon u_{xx} + \left(1 - \frac{\varepsilon}{2}\right) u_x + \frac{1}{4} \left(1 - \frac{1}{4}\varepsilon\right) u = e^{-x/4} \text{ on } [0, 1] \quad u(0) = u(1) = 0$$

$$u(x) = \exp^{\frac{-x}{4}} \left( x - \frac{\exp^{-\frac{1-x}{\varepsilon}} - \exp^{-\frac{1}{\varepsilon}}}{1 - \exp^{-\frac{1}{\varepsilon}}} \right)$$



# Numerical results: Convection Equation



# Convergence theory for PINNs

[Shin, Darbon and Kaniardarkis, 2020], [Jiao, Lai, Lo, Wang, Yang, 2024]

- PINN error is a combination of approximation error and training/optimization error
- Show that a PINN (depth  $L$  width  $W$ ) can be constructed with low approximation error which reduces as the complexity of the PINN increases.
- Hope that the optimization error can be reduced to acceptable levels.
- Try things out on simple problems


## BUT

- Non convex (no guarantee of uniqueness or convergence of training)
- Badly conditioned
- PINNs are nonlinear function approximators. No easy bound on the solution error.
- Solutions can sometimes have no connection to reality!

Compare with the PINN formulation with collocation

- PINN is a nonlinear approximation
- It may have better expressivity
- The collocation equations become an **ill conditioned** optimisation problem
- We may not find a good optimum due to the **lack of convexity and the ill conditioning**
- Pre conditioning and the right choice of optimiser are essential

# Detailed PINN Convergence Theory 1



## Detailed PINN Convergence Theory 2

Underlying solution:  $u(x)$

$H_n$ : Class of functions approximated by NN with  $n$  degrees of freedom

- $u_{\theta,m}$  : Approximation found in practice after some training. Is less accurate than
- $u_{\theta_m^*}$  : Approximation obtained by perfect optimisation with finite collocation points. Is less accurate than
- $u_{\theta^*,\infty}$  : Approximation obtained by perfect optimisation with infinite collocation points.

Measured error is  $\|u_{\theta,m} - u\|$



- $\|u_{\theta,m} - u_{\theta^*,m}\|$ : Optimisation error  $\mathcal{E}_O$ : Hard to control and depends on the optimiser and the initialisation:
- $\|u_{\theta^*,m} - u_{\theta^*,\infty}\|$ : Estimation error  $\mathcal{E}_E$ : Main results on this
- $\|u_{\theta^*,\infty} - u^*\|$ : Approximation error  $\mathcal{E}_A$  :

# Results on a mildly singular BVP

Compare standard collocation and a PINN for the ODE using ADAM

$$-\epsilon^2 u_{xx} + u = 1, \quad u(0) = u(1) = 0, \quad 0 < \epsilon \ll 1$$


This has a boundary layer at  $x = 0, x = 1$  of width  $\epsilon$ .



PINN

Collocation

# Using the LBFGS optimiser




# How and why do PINNs struggle to train .. and how can we fix this?

A subject of intensive research! See [Wang et. al.] for a summary:  
Observe

- PINNs train especially badly when they have high frequency features
- Loss function has multi-scale interactions, leading to ill-conditioning and stiffness in the gradient flow dynamics and a stringent stability requirement on the loss rate
- Standard networks show **ill-conditioning** and **spectral bias** and cannot learn high frequencies. This has been seen for example in **weather forecasting**

They advocate the use of **Neural Tangent Kernels** in the optimisation process. [Kiyani et. al.] extend this with the development of better optimisers eg. **Broyden methods**, for PINNs

Loss landscape is highly non-convex, [Kiyani et. al.]



Also Issues with **conditioning** as in the last lecture on approximation.  
Can fix for shallow networks by **preconditioning**

# Deep Ritz Method: Motivation the Calculus of variations

Seek to solve PDE problems eg. of the form

$$-\epsilon^2 u_{xx} + u = 1, \quad u(0) = u(1) = 0, \quad \Omega = [0, 1] \quad (*)$$

A PINN minimises the PDE residual as the loss function

**Instead use ideas from the calculus of variations to find a loss function**

**Functional minimisers:** Consider the functional

$$F(u) = \int_{\Omega} \frac{\epsilon^2 u_x^2}{2} + \frac{u^2}{2} - u \, dx.$$

**Then  $F(u)$  is minimised when  $u$  is the (unique) solution of the PDE (\*)**

# Approach 1: Finite Element Methodology

Express  $u(x)$  as a Galerkin approximation:

$$u(x) \approx U(x) = \sum_{i=0}^N U_i(t) \phi_i(x)$$

with  $\phi_i(x)$  **Not** globally differentiable *locally supported, piece-wise polynomial spline functions*

$U$  is in the linear space  $\mathcal{V}$  spanned by the functions  $\phi_i(x)$ .

- Minimise  $F$  over all such functions  $U$ .
- If the PDE is **linear**, this leads to a **linear** system for the coefficients  $U_i$  of the approximation
- Linear system is **symmetric** and **well conditioned**
- Solve this system using (for example) a conjugate-gradient method.

## Features of the FE method


- Unique solution (if the PDE is linear)
- Have guaranteed error estimates for all  $N$ , due to **Céa's lemma**, of the form

$$\|u - U\|_{H^1} < C(u)N^{-\alpha}$$

- Can reduce  $C(u)$  and increase  $\alpha$  using an adaptive approach.
- But .. Awkward in higher dimensions!

# Meshes

Traditional PDE computations using **Finite Element Methods** use a **computational mesh**  $\tau$  comprising mesh points and a mesh topology with  $\phi_i(x)$  defined over the mesh



**Meshes can be tricky to construct!** Although ML methods can help a lot, see [B, Maierhofer, Rowbottom et. al, ICML 2025]

## Approach 2: The Deep Ritz Method [Weinan E and Bing Yu, (2017)]

### Lovely Idea!!!

Let  $y(x)$  be the output of a parametrised NN

$$y(x) = \mathcal{N}N(x)$$

with  $\mathcal{Y}$  the (nonlinear) set of functions parametrised by  $\theta$

Then set

$$y^* = \operatorname{argmin}_{\mathcal{Y}} F.$$

Allowing for the boundary conditions

# Deep Ritz method for the Poisson equation

The **Deep Ritz Method** (DRM) seeks the solution  $u$  satisfying

$$y = \arg \min_{v \in H} \mathcal{F}(v),$$

where  $H$  is the set of admissible functions (trial functions) and

$$\mathcal{I}(v) = \int_{\Omega} \left( \frac{1}{2} |\nabla v(\mathbf{x})|^2 - f(\mathbf{x})v(\mathbf{x}) \right) d\mathbf{x} + \beta \int_{\partial\Omega} (v(\mathbf{x}) - u_D)^2 d\mathbf{s}$$

The Deep Ritz method is based of the following assumptions:

- $y(\mathbf{x})$  is NN based approximation of  $u$
- A numerical quadrature rule for the functional using chosen quadrature points eg. random, optimal
- An algorithm for solving the optimization problem eg. SGD on random quadrature points

## DRM vs.Finite Element

The DRM is superficially similar to the Finite Element method but has crucial differences as **the NN approximating subspace is nonlinear**, and the problem of finding the weights is **non-convex and badly conditioned**.

## FE: linear

- Limited expressivity, reduced accuracy
- Adaptive only with effort
- Not equivariant
- Need a complex mesh data structure
- Convex with guarantees of uniqueness for many problems (and direct calculation using linear algebra)
- Well conditioned
- Work on saddle-point problems (eg. most problems)
- Good (a-priori and a-posteriori) guaranteed error bounds :

**Cea's Lemma:** Bounds solution error by interpolation error on the FE space





## DRM:nonlinear

- Very expressive (potential high accuracy for a small number of degrees of freedom)
- Self adaptive
- Equivariant
- Don't need a complex mesh data structure
- Ill conditioned
- Don't work on saddle-point problems eg. most problems for example:

$$u_{xx} + u = 1, \quad u_x x + u^3 = 0.$$

- Don't have Céas Lemma


$$\text{Example 1: } -\epsilon u'' + u = 1, \quad u(0) = u(1) = 0.$$



## Example 2: Poisson equation in 2D:

DRM method works well for small DOF

dG Finite Element Method is **much** better for more DOF



This convergence pattern is seen in many other examples

# Summary

- PINNS and DRMS both show promise as a quick way of solving PDEs  
but have only really been tested on quite simple problems so far
- PINS not (yet) competitive with FE in like-for-like comparisons
- PINNs need careful meta-parameter tuning and preconditioning to work well
- Long way to go before we understand PINNS orn DRMS completey and have a satisfactory convergence theory for them in the general case.
- Lots of great stuff to do!